

dsc-project

November 30, 2016

```
In [1]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics.cluster import v_measure_score
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from time import time

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import StratifiedShuffleSplit

import warnings
warnings.filterwarnings('ignore')

In [2]: from sklearn.utils import resample
from sklearn import preprocessing

X = np.loadtxt('../data.txt')
Y = np.loadtxt('../label.txt').astype(int)

X = np.nan_to_num(X)

Xp, Yp = resample(X, Y, n_samples = 5000)

Xp_scaled = preprocessing.scale(Xp)

In [3]: def trial(X, Y, method, name, scoring='f1'):
    print (name)
    start_time = time()
    method.fit(X, Y)
    y_pred = method.predict(X)
```

```

print(classification_report(Y, y_pred))
print ("\tDone in %.2f s" % (time() - start_time))
print ("Cross validation...")
start_time = time()
skf = StratifiedKFold(n_splits=10, shuffle=True)
scores = cross_val_score(method, X, Y, cv=skf, scoring=scoring, n_jobs=-1)
accu = cross_val_score(method, X, Y, cv=skf, n_jobs=-1)
print ("\tAccuracy: %.2f (+/- %.2f)" % (accu.mean(), accu.std() * 2))
print ("\tF1 score: %.2f (+/- %.2f)" % (scores.mean(), scores.std() * 2))
print ("\tDone in %.2f s" % (time() - start_time))

```

0.1 Visualization

0.1.1 PCA and Kernel PCA

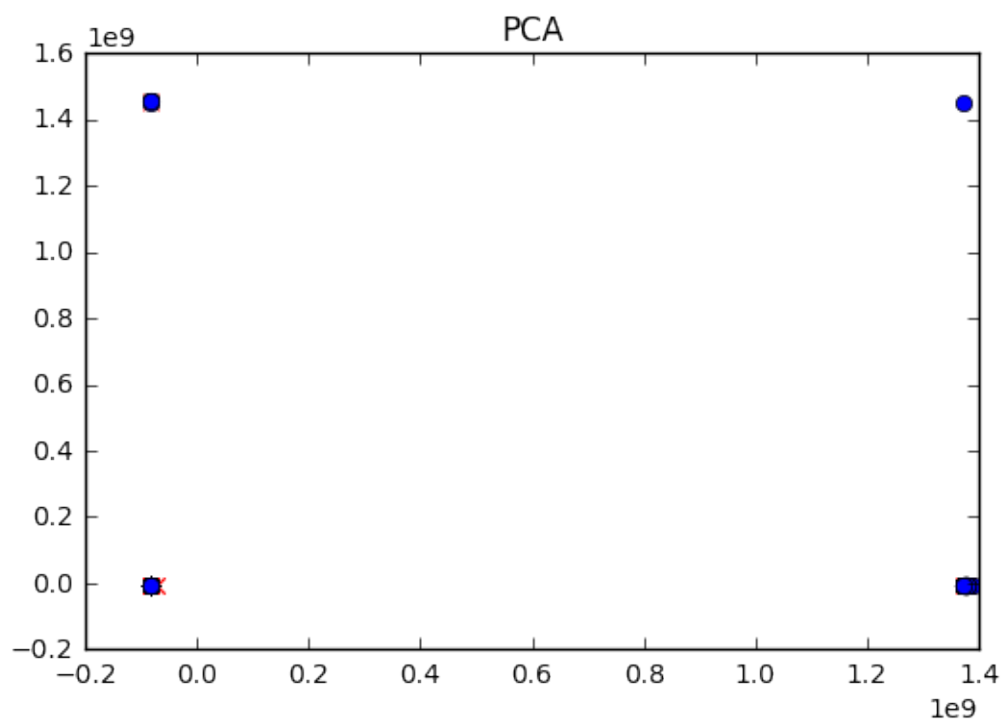

In [4]: `from sklearn.decomposition import PCA, KernelPCA`

```

kpca = KernelPCA(n_components = 2, kernel="rbf")
X_kpca = kpca.fit_transform(Xp)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(Xp)
reds = Yp == 0
blues = Yp == 1

plt.plot(X_kpca[reds, 0], X_kpca[reds, 1], "rx")
plt.plot(X_kpca[blues, 0], X_kpca[blues, 1], "bo")
plt.title('Kernel PCA')
plt.show()
plt.plot(X_pca[reds, 0], X_pca[reds, 1], "rx")
plt.plot(X_pca[blues, 0], X_pca[blues, 1], "bo")
plt.title('PCA')
plt.show()

```



0.2 random guess

```
In [5]: y_pred = np.random.randint(0, 2, Y.shape[0])
        scores = np.array([f1_score(Y, np.random.randint(0, 2, Y.shape[0])) for i in range(100)])
        accu = np.array([accuracy_score(Y, np.random.randint(0, 2, Y.shape[0])) for i in range(100)])
        print ("Accuracy: %0.2f (+/- %0.2f)" % (accu.mean(), accu.std() * 2))
        print ("F1 score: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

Accuracy: 0.50 (+/- 0.00)

F1 score: 0.40 (+/- 0.00)

0.3 Linear methods

```
In [6]: from sklearn.linear_model import SGDClassifier, Perceptron
        from sklearn.linear_model import PassiveAggressiveClassifier
        from sklearn.linear_model import LogisticRegression
```

```
In [7]: trial(X, Y, Perceptron(), 'Perceptron')
```

Perceptron

	precision	recall	f1-score	support
0	0.71	0.99	0.83	324701
1	0.93	0.17	0.28	161347
avg / total	0.78	0.72	0.64	486048

Done in 0.99 s

Cross validation...

Accuracy: 0.60 (+/- 0.35)

F1 score: 0.37 (+/- 0.21)

Done in 15.18 s

```
In [8]: trial(X, Y, LogisticRegression(), 'Log Reg')
```

Log Reg

	precision	recall	f1-score	support
0	0.71	0.99	0.83	324701
1	0.93	0.17	0.28	161347
avg / total	0.78	0.72	0.64	486048

Done in 2.40 s

Cross validation...

Accuracy: 0.72 (+/- 0.00)

F1 score: 0.28 (+/- 0.01)

Done in 68.48 s

```
In [9]: trial(X, Y, SGDClassifier(), 'SGD Clasifier')
```

SGD Clasifier

	precision	recall	f1-score	support
0	0.58	0.00	0.00	324701
1	0.33	1.00	0.50	161347
avg / total	0.50	0.33	0.17	486048

Done in 1.48 s

Cross validation...

Accuracy: 0.68 (+/- 0.23)

F1 score: 0.30 (+/- 0.25)

Done in 15.49 s

```
In [10]: trial(X, Y, SGDClassifier(loss="log", penalty="l2"), 'SGD with log')
```

SGD with log

	precision	recall	f1-score	support
0	0.00	0.00	0.00	324701
1	0.33	1.00	0.50	161347
avg / total	0.11	0.33	0.17	486048

Done in 1.05 s

Cross validation...

Accuracy: 0.60 (+/- 0.35)

F1 score: 0.39 (+/- 0.22)

Done in 15.78 s

```
In [11]: pac = PassiveAggressiveClassifier(random_state=9,
                                             class_weight='balanced',
                                             n_jobs=-1,
                                             n_iter=9)
         trial(X, Y, pac, 'Passive Aggresive Clasifier')
```

Passive Aggresive Clasifier

	precision	recall	f1-score	support
0	0.71	0.99	0.83	324701
1	0.93	0.17	0.28	161347
avg / total	0.78	0.72	0.64	486048

Done in 1.87 s

Cross validation...

Accuracy: 0.56 (+/- 0.38)

F1 score: 0.37 (+/- 0.21)

Done in 23.27 s

0.4 Non - Linear models

```
In [12]: from sklearn import svm
import warnings
warnings.filterwarnings('ignore')
```

```
In [13]: # Subsample
trial(Xp, Yp, svm.SVC(), 'SVC')
```

SVC

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3343
1	1.00	1.00	1.00	1657
avg / total	1.00	1.00	1.00	5000

Done in 3.95 s

Cross validation...

Accuracy: 0.67 (+/- 0.01)

F1 score: 0.02 (+/- 0.02)

Done in 28.05 s

```
In [14]: # Subsample
trial(Xp, Yp, svm.NuSVC(gamma=1e9), 'Nu SVC')
```

Nu SVC

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3343
1	1.00	1.00	1.00	1657
avg / total	1.00	1.00	1.00	5000

Done in 3.78 s

Cross validation...

Accuracy: 0.67 (+/- 0.00)

F1 score: 0.01 (+/- 0.02)

Done in 29.39 s

0.5 Non linear Transformations

```
In [15]: from sklearn.ensemble import RandomTreesEmbedding, ExtraTreesClassifier

# use RandomTreesEmbedding to transform data
hasher = RandomTreesEmbedding(n_jobs=-1)
X_randomTrees = hasher.fit_transform(X)
```

```
In [16]: trial(X_randomTrees, Y, LogisticRegression(), 'Random Trees Embedding')
```

Log Reg

	precision	recall	f1-score	support
0	0.71	0.99	0.83	324701
1	0.92	0.17	0.28	161347
avg / total	0.78	0.72	0.65	486048

Done in 9.45 s

Cross validation...

Accuracy: 0.72 (+/- 0.00)

F1 score: 0.28 (+/- 0.01)

Done in 141.27 s

```
In [17]: trees = ExtraTreesClassifier()
        trial(Xp, Yp, trees, 'trees')
```

trees

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3343
1	1.00	1.00	1.00	1657
avg / total	1.00	1.00	1.00	5000

Done in 0.10 s

Cross validation...

Accuracy: 0.70 (+/- 0.03)

F1 score: 0.41 (+/- 0.06)

Done in 1.25 s

```
In [18]: from sklearn.naive_bayes import BernoulliNB
        nb = BernoulliNB()
        trial(X_randomTrees, Y, nb, 'Naive bayes')
        trial(X, Y, nb, 'Naive bayes')
```

Naive bayes

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.71	0.96	0.82	324701
1	0.73	0.19	0.31	161347
avg / total	0.72	0.71	0.65	486048

Done in 0.35 s
 Cross validation...
 Accuracy: 0.71 (+/- 0.00)
 F1 score: 0.31 (+/- 0.01)
 Done in 4.56 s

Naive bayes

	precision	recall	f1-score	support
0	0.72	0.86	0.78	324701
1	0.54	0.33	0.41	161347
avg / total	0.66	0.68	0.66	486048

Done in 1.18 s
 Cross validation...
 Accuracy: 0.68 (+/- 0.00)
 F1 score: 0.41 (+/- 0.01)
 Done in 41.89 s

In [19]: `from sklearn.kernel_approximation import RBFSampler`

In [20]: `rbf_feature = RBFSampler(gamma=1, random_state=1)`
`X_rbf = rbf_feature.fit_transform(Xp)`
`trial(X_rbf, Yp, trees, 'RBF transformation')`

RBF transformation

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3343
1	1.00	1.00	1.00	1657
avg / total	1.00	1.00	1.00	5000

Done in 0.29 s
 Cross validation...
 Accuracy: 0.64 (+/- 0.03)
 F1 score: 0.16 (+/- 0.08)
 Done in 1.54 s

0.6 Manifold learning

```
In [21]: from sklearn import neighbors
         nnc = neighbors.KNeighborsClassifier(n_neighbors = 1,
                                             weights='uniform',
                                             algorithm='kd_tree')

         trial(Xp, Yp, nnc, 'K Neighbors')
```

Nearest neighbors

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3343
1	1.00	1.00	1.00	1657
avg / total	1.00	1.00	1.00	5000

Done in 0.83 s

Cross validation...

Accuracy: 0.61 (+/- 0.04)

F1 score: 0.41 (+/- 0.05)

Done in 1.58 s

```
In [22]: nnc = neighbors.KNeighborsClassifier(n_neighbors = 4,
                                             weights='distance',
                                             algorithm='auto')

         trial(Xp, Yp, nnc, 'Nearest neighbors')
```

Nearest neighbors

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3343
1	1.00	1.00	1.00	1657
avg / total	1.00	1.00	1.00	5000

Done in 0.69 s

Cross validation...

Accuracy: 0.63 (+/- 0.03)

F1 score: 0.39 (+/- 0.04)

Done in 1.44 s

```
In [23]: from sklearn.neighbors import RadiusNeighborsClassifier

         rnc = RadiusNeighborsClassifier(radius=100)
         trial(Xp_scaled, Yp, rnc, 'Radius neighbors classiflier')
```

Radius neighbors classiflier

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.67	1.00	0.80	3343
1	0.00	0.00	0.00	1657
avg / total	0.45	0.67	0.54	5000

Done in 15.36 s
Cross validation...
Accuracy: 0.67 (+/- 0.00)
F1 score: 0.00 (+/- 0.00)
Done in 12.97 s

0.7 Multi layer peceptron

```
In [24]: from sklearn.neural_network import MLPClassifier
```

```
mlp = MLPClassifier(solver='lbfgs', alpha=1e-5, activation='relu',
                    hidden_layer_sizes=(100,33), random_state=10, tol=1e-9,
                    max_iter=400)
```

```
trial(Xp, Yp, mlp, 'Multi layer perceptron')
```

Multi layer perceptron				
	precision	recall	f1-score	support
0	0.00	0.00	0.00	3343
1	0.33	1.00	0.50	1657
avg / total	0.11	0.33	0.16	5000

Done in 0.80 s
Cross validation...
Accuracy: 0.33 (+/- 0.00)
F1 score: 0.50 (+/- 0.00)
Done in 11.18 s

```
In [25]: parameters = {'solver': ['lbfgs'],
                        'alpha': 10.0 * -np.arange(-1, 7),
                        'hidden_layer_sizes': [(100,33)],
                        'activation': ['identity', 'logistic', 'tanh', 'relu'],
                        'max_iter' : [400],
                        'tol': [1e-5, 1e-9],
                        'learning_rate': ['constant', 'invscaling', 'adaptive']}

mlp = MLPClassifier()
# cv = StratifiedShuffleSplit(n_splits=10)
```

```

skf = StratifiedKFold(n_splits=10, shuffle=True)
clf = GridSearchCV(mlp, parameters, cv=skf, n_jobs=-1, scoring='f1')
start_time = time()
# clf.fit(Xp, Yp)
#
# The best parameters are {'tol': 1e-05, 'activation': 'identity', 'alpha': 1.0}
# Done in 11507.70 s

print("The best parameters are %s with a score of %0.2f"
      % (clf.best_params_, clf.best_score_))
print ("\tDone in %0.2f s" % (time() - start_time))

```

The best parameters are {'tol': 1e-05, 'activation': 'identity', 'alpha': 1.0, 'learning_rate': 0.001} Done in 11507.70 s

```

In [26]: mlp = MLPClassifier(solver='lbfgs', alpha=1e-6,
                           hidden_layer_sizes=(100,33),
                           learning_rate='invscaling',
                           tol = 1e-5)

```

```

trial(Xp, Yp, mlp, 'Multi layer perceptron')

```

	precision	recall	f1-score	support
0	0.67	1.00	0.80	3343
1	0.43	0.00	0.00	1657
avg / total	0.59	0.67	0.54	5000

```

Done in 0.60 s
Cross validation...
Accuracy: 0.67 (+/- 0.00)
F1 score: 0.19 (+/- 0.34)
Done in 12.21 s

```

```

In [27]: parameters = {'n_neighbors': [1, 2, 4, 10, 20],
                      'weights': ['uniform', 'distance']}

ncc = neighbors.KNeighborsClassifier()
# cv = StratifiedShuffleSplit(n_splits=10)
skf = StratifiedKFold(n_splits=10, shuffle=True)
clf = GridSearchCV(ncc, parameters, cv=skf, n_jobs=-1, scoring='f1')
start_time = time()
print ("start")
clf.fit(Xp, Yp)
print("The best parameters are %s with a score of %0.2f"

```

```

        % (clf.best_params_, clf.best_score_)
    print ("\tDone in %.2f s" % (time() - start_time))

```

start

The best parameters are {'weights': 'uniform', 'n_neighbors': 1} with a score of 0.
Done in 40.58 s

```

In [41]: parameters = {'n_estimators': [1024],
                        'min_samples_split': [256],
                        'class_weight': ['balanced', None],
                        'max_depth': [None],
                        'max_features': [5, None, 'sqrt']}

```

```

trees = ExtraTreesClassifier()
skf = StratifiedKFold(n_splits=10, shuffle=True)
clf = GridSearchCV(trees, parameters, cv=skf, n_jobs=-1, scoring='f1')
start_time = time()
'''start
The best parameters are {'class_weight': 'balanced_subsample', 'max_depth': 10}
Done in 1698.70 s
'''

```

```

print ("start")
clf.fit(Xp, Yp)
print ("The best parameters are %s with a score of %0.2f"
      % (clf.best_params_, clf.best_score_))
print ("\tDone in %.2f s" % (time() - start_time))

```

start

The best parameters are {'class_weight': 'balanced', 'max_features': None, 'max_depth': 10}
Done in 293.03 s

```

In [36]: trees = ExtraTreesClassifier(class_weight='balanced',
                                     n_estimators=1024,
                                     min_samples_split=256)

```

```

trial(Xp, Yp, trees, 'Extra tree classifier optimized')

```

Extra tree classifier optimized

	precision	recall	f1-score	support
0	0.81	0.81	0.81	3343
1	0.62	0.62	0.62	1657
avg / total	0.75	0.75	0.75	5000

```
    Done in 9.27 s
Cross validation...
    Accuracy: 0.68 (+/- 0.05)
    F1 score: 0.51 (+/- 0.05)
    Done in 72.17 s
```